

spar: Sparse Projected Averaged Regression in R

Roman Parzer 
TU Wien

Laura Vana Gür 
TU Wien

Peter Filzmoser 
TU Wien

Abstract

The **spar** package for R builds ensembles of predictive generalized linear models for high-dimensional data. It employs an algorithm that combines variable screening and random projection techniques to address the computational challenges of large predictor sets. This package makes this modeling approach more accessible, as variants of the algorithm have demonstrated competitive predictive performance compared to state-of-the-art methods, while maintaining low computational costs.

Designed with extensibility in mind, **spar** implements screening and random projection techniques as S3 classes with user-friendly constructor functions. This modular design allows users to seamlessly integrate and develop new procedures, making the package a versatile tool for high-dimensional applications.

Keywords: random projection, variable screening, ensemble learning, R.

1. Introduction

The **spar** package for R ([R Core Team 2024](#)) offers functionality for estimating generalized linear models (GLMs) in high-dimensional settings, where the number of predictors p significantly exceeds the number of observations n i.e., $p > n$ or even $p \gg n$. To address the challenges of high dimensionality, the package implements an algorithm which integrates variable screening methods with random projection techniques to effectively reduce the predictor space.

Random projection is a computationally-efficient method which linearly maps a set of points in high dimensions into a much lower-dimensional space. Several packages in R provide functionality for random projections. For instance, package **RandPro** ([Aghila and Siddharth 2020](#); [Siddharth and Aghila 2020](#)) allows a Gaussian random matrix (i.e, where entries are simulated from a standard normal), a sparse matrix ([Achlioptas 2003](#); [Li, Hastie, and Church 2006](#)) or a matrix generated using the equal probability distribution with the elements $\{-1, 1\}$ to be applied to the predictor matrix before employing one of k nearest neighbors, support vector machine or naive Bayes classifier on the projected features. Package **SPCAvRP** ([Gataric, Wang, and Samworth 2019](#)) implements sparse principal component analysis, based on the aggregation of eigenvector information from “carefully-selected” axis-aligned random projections of the sample covariance matrix. Additionally, package **RPEnsembleR** ([Cannings and](#)

Samworth 2021) implements a similar idea when building the ensemble of classifiers: for each classifier in the ensemble, a collection of (Gaussian, axis-aligned projections, or Haar) random projection matrices is generated, and the one that minimizes a risk measure for classification on a test set is selected. For Python (van Rossum *et al.* 2011) the `sklearn.random_projection` module (Pedregosa, Varoquaux, Gramfort, Michel, Thirion, Grisel, Blondel, Prettenhofer, Weiss, Dubourg *et al.* 2011) implements two types of unstructured random matrices, namely Gaussian random matrix and sparse random matrix.

Random projection can suffer from noise accumulation for very large p , as too many irrelevant predictors are being considered for prediction purposes (Mukhopadhyay and Dunson 2020). Therefore, screening out irrelevant variables before performing the random projection is advisable in order to tackle this issue, while also reducing computational costs. The screening procedure involves computing for each predictor variable a measure of its *utility* to the response variable and then selecting a small subset of variables based on this measure for subsequent analyses. The screening can be performed by selecting a predetermined number of best variables or in a probabilistic fashion, by randomly sampling covariates for inclusion in the model based on probabilities proportional to the utility measure (as opposed to random subspace sampling employed in, e.g., random forests). The R landscape for variable screening techniques is very rich. An overview of some notable packages on the Comprehensive R Archive Network (CRAN) includes the following packages. Package **SIS** (Saldana and Feng 2018), which implements the (iterative) sure independence screening procedure and its extensions, as detailed in Fan and Lv (2007), Fan and Song (2010), Fan, Feng, and Wu (2010). This package also provides functionality for estimating a penalized generalized linear model or a cox regression model for the variables selected by the screening procedure. Package **VariableScreening** (Li, Huang, and Dziak 2022) offers screening methods for independent and identically distributed (iid) data, varying-coefficient models, and longitudinal data and includes techniques such as sure independent ranking and screening (SIRS), which ranks the predictors by their correlation with the rank-ordered response, or distance correlation sure independence screening (DC-SIS), a non-parametric extension of the correlation coefficient. Package **MFSIS** (Cheng, Wang, Zhu, Zhong, and Zhou 2024) provides a collection of model-free screening techniques including SIRS, DC-SIS, the fused Kolmogorov filter (Mai and Zou 2015) the projection correlation method using knock-off features (Wanjin Liu and Li 2022), among others. Additional packages implement specific procedures but their review is beyond the scope of the current paper.

Although combining variable screening with random projection effectively reduces the predictor set and computational costs, the variability introduced by random sampling – both in projections and screening indices – can be mitigated by averaging the results from multiple iterations (Thanei, Heinze, and Meinshausen 2017). To address these points, **spar** builds an ensemble of GLMs in the spirit of Mukhopadhyay and Dunson (2020) and Parzer, Filzmoser, and Vana-Gür (2024a) where, in each model of the ensemble, i) variables are first screened based on a screening coefficient, ii) the selected variables are then projected to a lower dimensional space, iii) a GLM is estimated using the projected predictors. Finally, additional sparsity in the coefficients of the original variables can be introduced through a thresholding parameter, which together with the number of models in the ensemble can be chosen using a validation set or via cross-validation. The final coefficients are then obtained by averaging over the marginal models in the ensemble. This algorithm performs sparse projected averaged regression (SPAR) for both discrete and continuous data in the GLM framework in a com-

putationally efficient way. Different variants of the algorithm have been shown to perform well in terms of prediction power on a variety of data sets (see [Mukhopadhyay and Dunson 2020](#)). In particular, [Parzer et al. \(2024a\)](#) show that when paired with a carefully constructed data-driven random projection the algorithm performs superiorly in terms of predictions and variable ranking in settings exhibiting different degrees of sparsity in the coefficients.

A variety of screening coefficients are provided as well as several procedures for generating random projection matrices. These procedures can be classified into data-agnostic and data-driven, where the former, unlike the latter, does not incorporate information from the data in the construction of the matrices. However, the package is designed with flexibility in mind, providing users with a versatile framework to extend the implemented screening and projection techniques with their own custom procedures. This is facilitated by leveraging R’s S3 classes, making the process convenient and user-friendly. Therefore, users can seamlessly integrate various techniques by either writing their own procedures or leveraging the existing R packages.

The package provides methods such as `plot`, `predict`, `coef`, `print`, which allow users to more easily interact with the model output and analyze the results. The GLM framework, especially when combined with random projections which preserve information on the original coefficients (such as the one in [Parzer et al. 2024a](#)), facilitates interpretability of the model output, allowing users to understand variable effects.

While **spar** offers the first implementation of the described algorithm and, to the best of our knowledge, no other package offers the same functionality for GLMs, few other R packages focus on building ensembles where the dimensionality of the predictors is reduced. Most notably, package **RPEnsemble** ([Cannings and Samworth 2021](#)) implements the procedure in [Cannings and Samworth \(2017\)](#) where “carefully-selected” random projections are used for projecting the predictors before they are employed in a classifier such as k -nearest neighbor, linear or quadratic discriminant analysis. On the other hand, package **RaSEn** ([Tian and Feng 2021](#)) implements the RaSE algorithm for ensemble classification and regression problems, where random subspaces are generated and the optimal one is chosen to train a weak learner on the basis of some criterion.

The rest of the paper is organized as follows: Section 2 provides the methodological details of the implemented algorithm. The package is described in Section 3 and Section 4 exemplifies how a new screening coefficient and a new random projection can be integrated in the package. Section 5 contains two examples of employing the package on real data sets. Section 6 concludes.

2. Methods

The package implements a procedure for building an ensemble of GLMs where we employ screening and random projection to the predictor matrix pre-model estimation for the purpose of dimensionality reduction.

Throughout the section we assume to observe high-dimensional data $\{(x_i, y_i)\}_{i=1}^n$, where $x_i \in \mathbb{R}^p$ is a predictor vector and $y_i \in \mathbb{R}$ is the response, with $p \gg n$. The predictor vectors are collected in the rows of the predictor matrix $X \in \mathbb{R}^{n \times p}$.

2.1. Variable screening

The field of variable screening is an active area of research, with numerous methods developed to address different data settings. While a comprehensive review is beyond the scope of this paper, this section provides a concise and selective overview of key approaches for variable screening in high-dimensional settings.

In high-dimensional modeling, the goal of variable screening is to reduce the predictor set by selecting a small subset of variables with a strong *utility* to the response variable. This initial selection enables more efficient downstream analyses by discarding less relevant predictors early in the modeling process, thus reducing computational costs and potential noise accumulation stemming from irrelevant variables (see e.g., [Mukhopadhyay and Dunson 2020](#)).

A classic approach is sure independence screening (SIS), proposed by [Fan and Lv \(2007\)](#), which uses the vector of marginal empirical correlations $\hat{\omega} = (\omega_1, \dots, \omega_p)^\top \in \mathbb{R}^p$, $\omega_j = \text{Cor}(X_j, y)$, where y is the $(n \times 1)$ vector of responses and X_j is the j -th column of the matrix of predictors, to screen predictors in a linear regression setting by selecting the variable set $\mathcal{A}_\gamma = \{j \in [p] : |\omega_j| > \gamma\}$ depending on a threshold $\gamma > 0$, where $[p] = \{1, \dots, p\}$. Under certain technical conditions, this screening coefficient has the *sure screening property* $\mathbb{P}(\mathcal{A} \subset \mathcal{A}_{\gamma_n}) \rightarrow 1$ for $n \rightarrow \infty$, where $\mathcal{A} = \{j \in [p] : \beta_j \neq 0\}$ is the set of truly active variables. Extensions to SIS include modifications for GLMs ([Fan and Song 2010](#)), where screening is performed based on the log-likelihood $\ell(\cdot)$ or the slope coefficient of the GLM containing only X_j as a predictor: $\hat{\omega}_j =: \arg\min_{\beta_j \in \mathbb{R}} \min_{\beta_0 \in \mathbb{R}} \sum_{i=1}^n -\ell(\beta_j, \beta_0; y_i, x_{ij})$, where x_{ij} is the j -th entry of the vector x_i .

However, both mentioned approaches face limitations related to the required technical conditions which can rule out practically possible scenarios where an important variable is marginally uncorrelated to the response due to their multicollinearity. To tackle these issues, [Fan, Samworth, and Wu \(2009\)](#) propose to use an iterative procedure where SIS is applied subsequently on the residuals of the model estimated in a previous step. Additionally, in a linear regression setting, [Wang and Leng \(2016\)](#) propose employing the ridge estimator when the penalty term converges to zero while [Cho and Fryzlewicz \(2012\)](#) propose using the tilted correlation, i.e., the correlation of a tilted version of X_j with y where the effect of other variables is reduced. For discrete outcomes, joint feature screening ([Xu and Chen 2014](#)) has been proposed.

In order to tackle potential model misspecification, a rich stream of literature focuses on developing semi- or non-parametric alternatives to SIS. For linear regression, approaches include using the ranked correlation ([Zhu, Li, Li, and Zhu 2011](#)), (conditional) distance correlation ([Li, Zhong, and Zhu 2012](#); [Wang, Pan, Hu, Tian, and Zhang 2015](#)) or quantile correlation ([Ma and Zhang 2016](#)). For GLMs, [Fan, Feng, and Song \(2011\)](#) extend [Fan and Song \(2010\)](#) by fitting a generalized additive model with B-splines. Further extensions for discrete (or categorical) outcomes include the fused Kolmogorov filter ([Mai and Zou 2013](#)), the mean conditional variance, i.e., the expectation in X_j of the variance in the response of the conditional cumulative distribution function $\mathbb{P}(X \leq x|Y)$ ([Cui, Li, and Zhong 2015](#)). [Ke \(2023\)](#) propose a model free method where the contribution of each individual predictor is quantified marginally and conditionally in the presence of the control variables as well as the other candidates by reproducing-kernel-based R^2 and partial R^2 statistics.

Package **spar** allows the integration of such (advanced) screening techniques using a flexi-

ble framework, which in turn enables users to apply various screening methods tailored to their data characteristics in the algorithm generating the ensemble. This flexibility allows users to evaluate different strategies, ensuring that the most effective approach is chosen for the specific application at hand. Moreover, it incorporates probabilistic screening strategies, which can be particularly useful in ensembles, as they enhance the diversity of predictors across ensemble models. Instead of relying on a fixed threshold or number of predictors to be screened, predictors are sampled with probabilities proportional to their screening score (see Mukhopadhyay and Dunson 2020; Parzer *et al.* 2024a).

2.2. Random projection tools

Package **spar** has been designed to allow the incorporation of various random projection techniques, enabling users to tailor the procedure to their specific data needs. Below, we provide background information on random projection techniques and an overview of possible choices for building such random projection matrices.

The random projection method relies on the Johnson-Lindenstrauss (JL) lemma (Johnson and Lindenstrauss 1984), which asserts that for each set of points in p -dimensional Euclidean space collected in the rows of $X \in \mathbb{R}^{n \times p}$ there exists a linear map $\Phi \in \mathbb{R}^{m \times p}$ such that all pairwise distances are approximately preserved within a factor of $(1 \pm \epsilon)$ for $m \geq m_0 = \mathcal{O}(\epsilon^{-2} \log(n))$. Computationally, an attractive feature of the method for high-dimensional settings is that the bound does not depend on p .

The goal is to choose a random map Φ that satisfies the JL lemma with high probability given that it fulfills certain technical conditions. The literature focuses on constructing such matrices either by sampling them from some “appropriate” distribution, by inducing sparsity in the matrix and/or by employing specific fast constructs which lead to efficient matrix-vector multiplications. It turns out that the conditions are generally satisfied by nearly all sub-Gaussian distributions (Matoušek 2008). A common choice is the standard normal distribution $\Phi_{ij} \stackrel{iid}{\sim} N(0, 1)$ (Frankl and Maehara 1988) or a sparser version where $\Phi_{ij} \stackrel{iid}{\sim} N(0, 1/\sqrt{\psi})$ with probability ψ and 0 otherwise (Matoušek 2008). Another computationally simpler option is the Rademacher distribution where $\Phi_{ij} = \pm 1/\sqrt{\psi}$ with probability $\psi/2$ and zero otherwise for $0 < \psi \leq 1$, where Achlioptas (2003) shows results for $\psi = 1$ and $\psi = 1/3$ while Li *et al.* (2006) recommend using $\psi = 1/\sqrt{p}$ to obtain very sparse matrices.

Further approaches include using the Haar measure to generate random orthogonal matrices (Cannings and Samworth 2017) or a non-sub-Gaussian distribution like the standard Cauchy, proposed by Li *et al.* (2006) for preserving approximate ℓ_1 distances in settings where the data is high-dimensional, non-sparse, and heavy-tailed. Structured matrices, which allow for more efficient multiplication, have also been proposed (see e.g., Ailon and Chazelle 2009; Clarkson and Woodruff 2013).

The conventional random projections mentioned above are data-agnostic. However, recent work has proposed incorporating information from the data either to select the “best” data-agnostic random projection or to directly inform the random projection procedure. Cannings and Samworth (2017) rely on the former approach and build an ensemble classifier where the random projection matrix is chosen by selecting the one that minimizes the test error of the classification problem among a set of data-agnostic random projections.

On the other hand, [Parzer *et al.* \(2024a\)](#) propose to use a random projection matrix for GLMs which directly incorporates information about the relationship between the predictors and the response in the projection matrix, rather than a projection matrix which satisfies the JL lemma. [Parzer, Filzmoser, and Vana-Gür \(2024b\)](#) also provide in the linear regression a theoretical bound on the expected gain in prediction error in using a projection which incorporates information about the true regression coefficients compared to a conventional random projection. Motivated by this result, they propose to construct a projection matrix using the sparse embedding matrix of [Clarkson and Woodruff \(2013\)](#), where the random diagonal elements are replaced in practice by a ridge coefficient with a minimal λ penalty. This method has the advantage of approximately capturing the true regression coefficients in the span of the random projection, i.e., it ensures that the true regression coefficients can be recovered approximately after the projection.

Another data-driven approach to random projection for regression has been proposed by [Ryder, Karnin, and Liberty \(2019\)](#), who propose a data-informed random projection using an asymmetric transformation of the predictor matrix without using information of the response.

2.3. Generalized linear models

After we perform in each marginal model an initial screening step followed by a projection step, we assume that the reduced and projected set of predictors z_i together with the response arise from a GLM with the response having conditional density from a (reproductive) exponential dispersion family of the form

$$f(y_i|\theta_i, \phi) = \exp\left\{\frac{y_i\theta_i - b(\theta_i)}{a(\phi)} + c(y_i, \phi)\right\}, \quad g(\mathbb{E}[y_i|z_i]) = \gamma_0 + z_i^\top \gamma =: \eta_i,$$

where θ_i is the natural parameter, $a(\cdot) > 0$ and $c(\cdot)$ are specific real-valued functions determining different families, ϕ is a dispersion parameter, and $b(\cdot)$ is the log-partition function normalizing the density to integrate to one. If ϕ is known, we obtain densities in the natural exponential family for our responses. The responses are related to the m -dimensional reduced and projected predictors through the conditional mean, i.e., the conditional mean of y_i given z_i depends on a linear combination of the predictors through a (invertible) link function $g(\cdot)$, where $\gamma_0 \in \mathbb{R}$ is the intercept and $\gamma \in \mathbb{R}^m$ is a vector of regression coefficients for the m projected predictors.

Estimates for $\gamma_0 \in \mathbb{R}$ and $\gamma \in \mathbb{R}^m$ can be obtained by maximum likelihood if $m < n$. Even if $m < n$ (we recommend using a dimension m of at most $n/2$ – see [Section 2.4](#)), it might still be desirable to add a (e.g., small L_2) penalty to the likelihood of the marginal models – for the binomial family this can alleviate problems related to separation. Moreover, if outlier influence should be reduced, M - or trimmed estimators could be employed for obtaining a robust estimated of the regression coefficients (see e.g., [Cantoni and Ronchetti 2001](#)).

2.4. SPAR algorithm

We present the general algorithm for sparse projected averaged regression (SPAR) implemented in package **spar**.

1. Choose family with corresponding log-likelihood $\ell(\cdot)$ and link.

2. Standardize the $(n \times p)$ matrix of predictors X for all families and the vector of responses y for the Gaussian family by subtracting the sample column mean and dividing by the sample standard deviation.
3. Calculate screening coefficients $\hat{\omega}$.
4. For $k = 1, \dots, M$ models:
5. If $p > 2n$, screen $2n$ predictors based on the screening coefficient $\hat{\omega}$, which yields for model k the screening index set $I_k = \{j_1^k, \dots, j_{2n}^k\} \subset [p]$; if probabilistic screening should be employed draw the predictors sequentially without replacement using an initial vector of probabilities $p_j \propto |\hat{\omega}_j|$. Otherwise, select the $2n$ variables with the highest $|\hat{\omega}_j|$. If $p < 2n$, perform no screening and $I_k = \{1, \dots, p\}$.
6. project screened variables to a random dimension $m_k \sim \text{Unif}\{\log(p), \dots, n/2\}$ using **projection matrix** Φ_k to obtain $Z_k = X_{\cdot I_k} \Phi_k^\top \in \mathbb{R}^{n \times m_k}$, where $X_{\cdot I_k}$ contains the columns in X having a column index in I_k .
7. fit a **GLM**-type of model of y on Z_k to obtain estimated coefficients $\hat{\gamma}^k \in \mathbb{R}^{m_k}$ and $\hat{\beta}_{I_k}^k = \Phi_k^\top \hat{\gamma}^k$, $\hat{\beta}_{\bar{I}_k}^k = 0$.
8. For a given threshold $\nu > 0$, set all $\hat{\beta}_j^k$ with $|\hat{\beta}_j^k| < \nu$ to 0 for all j, k .
9. Choose M and ν via a validation set or cross-validation by repeating steps 1 to 4 and employing a loss function $K(M, \nu)$ on the test set

$$(M_{\text{best}}, \nu_{\text{best}}) = \text{argmin}_{M, \nu} K(M, \nu).$$

10. Combine models of the ensembles via the coefficients using **simple average** $\hat{\beta} = \sum_{k=1}^M \hat{\beta}^k / M$.
11. Output the estimated coefficients and predictions for the chosen M and ν .

3. Software

The package can be installed from GitHub

```
R> devtools::install_github("RomanParzer/SPAR")
```

and loaded by

```
R> library("spar")
```

In this section we rely for illustration purposes on the example data set from the package which contains $n = 200$ observations of a continuous response y and $p = 2000$ predictors x which can be used as a training data set and $n = 100$ observations to be used as a test set.

```
R> data("example_data", package = "spar")
R> str(example_data)
#> List of 7
#> $ x      : num [1:200, 1:2000] 1.8302 -0.4251 -1.3893 -0.0947 0.4304 ...
#> $ y      : num [1:200] -5.64 -23.63 -17.09 13.18 20.91 ...
#> $ xtest  : num [1:100, 1:2000] -0.166 -0.3729 0.0379 0.6774 0.2174 ...
#> $ ytest  : num [1:100] 10.61 -34.1 29.3 35.53 8.67 ...
#> $ mu     : num 1
#> $ beta   : num [1:2000] 1 -2 3 2 1 -3 2 3 1 -2 ...
#> $ sigma2 : num 83
```

This data set has been simulated from a linear regression model with $\sigma^2 = 83$, an intercept $\mu = 1$ and β coefficients with 100 non-zero entries, where the non-zero entries are uniformly sampled from $\{-3, -2, -1, 1, 2, 3\}$.

3.1. Main functions and their arguments

The two main functions for fitting the SPAR algorithm are:

```
spar(x, y, family = gaussian("identity"), model = NULL,
rp = NULL, screencoef = NULL,
xval = NULL, yval = NULL, nnu = 20, nus = NULL, nummods = c(20),
measure = c("deviance", "mse", "mae", "class", "1-auc"),
inds = NULL, RPMs = NULL, ...)
```

which implements the algorithm in Section 2.4 without cross-validation and returns an object of class ‘spar’, and

```
spar.cv(x, y, family = gaussian("identity"), model = NULL,
rp = NULL, screencoef = NULL,
nfolds = 10, nnu = 20, nus = NULL, nummods = c(20),
measure = c("deviance", "mse", "mae", "class", "1-auc"), ...)
```

which implements the cross-validated procedure and returns an object of class ‘spar.cv’.

The common arguments of these functions are:

- **x** an $n \times p$ numeric matrix of predictor variables where $n < p$,
- **y** numeric response vector of length n ,
- **family** object from `stats::family()`; defaults to `gaussian()`;
- **model** an object of class ‘sparmodel’ which specifies the model employed for each element of the ensemble.
- **rp** an object of class ‘randomprojection’, defaults to `rp_cw(data = TRUE)`;
- **screencoef** an object of class ‘screencoef’, defaults to `screen_glmnet()`;
- **nnu** is the number of threshold values ν which should be considered for thresholding; defaults to 20;
- **nus** is an optional vector of ν values to be considered for thresholding. If it is not provided, is defaults to a grid of **nnu** values. This grid is generated by including zero and **nnu**–1 quantiles of the absolute values of the estimated coefficients from the marginal models, chosen to be equally spaced on the probability scale .
- **nummods** is the number of models to be considered in the ensemble; defaults to 20. If a vector is provided, all combinations of **nus** and **nummods** are considered when choosing the optimal ν_{best} and M_{best} .
- **measure** specifies the measure $K(\nu, M)$ based on which the thresholding value ν_{opt} and the number of models **M** should be chosen on the validation set (for `spar()`) or in each of the folds (in `spar.cv()`). The default value for **measure** is "deviance",

which is available for all families. Other options are mean squared error "mse" or mean absolute error "mae" (between responses and predicted conditional means, for all families), "class" (misclassification error) and "1-auc" (one minus area under the ROC curve) both just for binomial family.

Furthermore, `spar()` has the specific arguments:

- `xval` and `yval` which are used as validation sets for choosing ν_{best} and M_{best} . If not provided, `x` and `y` will be employed.
- `inds` is an optional list of length `max(nummods)` containing column index-vectors corresponding to variables that should be kept after screening for each marginal model; dimensions need to fit those of the dimensions of the provided matrices in RPM.
- `RPMs` is an optional list of length `max(nummods)` which contains projection matrices to be used in each marginal model.

Function `spar.cv()` has the specific argument `nfolds` which is the number of folds to be used for cross-validation. It relies on a lightweight version of `spar()` as a workhorse, which is called for each fold. The random projections for each model are held fixed throughout the cross-validation to reduce the computational burden. However, it is possible to specify whether the data associated with the random projection (relevant for data-driven random projections) should be updated in each fold iteration with the corresponding training data. This is achieved by modifying elements of the 'randomprojection' object, which we will exemplify in Section 4.

3.2. Screening coefficients

The objects for creating screening coefficients are implemented as S3 classes 'screencoef'. These objects are created by several implemented `screen_*` functions, which take as arguments ... (to be saved as attributes) and `control` (a list of controls to be used in the main function for computing the screening coefficient).

Consider as an example function `screen_marglik()` which implements a screening procedure based on the coefficients of univariate GLMs:

```
R> screen_marglik
#> function(..., control = list()) {
#>   out <- list(name = name,
#>               generate_fun = generate_fun,
#>               control = control)
#>   attr <- list2(...)
#>   attributes(out) <- c(attributes(out), attr)
#>   if (is.null(attr(out, "type"))) {
#>     attr(out, "type") <- "prob"
#>   } else {
#>     stopifnot(
#>       "'type' must be either 'prob' or 'fixed'." =
#>       (attr(out, "type") == "prob" | attr(out, "type") == "fixed")
#>     )
#> }
```

```

#>   }
#>   if (is.null(attr(out, "reuse_in_rp"))) attr(out, "reuse_in_rp") <- FALSE
#>   class(out) <- c("screencoef")
#>   return(out)
#> }
#> <bytecode: 0x119895e78>
#> <environment: 0x119883040>

```

The function returns a ‘**screencoef**’ object which is a list containing three elements: a character **name**, a function **generate_fun** returning a p dimensional vector of the screening coefficients and a list of **control** arguments useful in **generate_fun**. More details on the construction of such objects is

As mentioned above, arguments related to the screening procedure can be passed to the **screen_*()** function through $\dots$, and will be saved as attributes of the ‘**screencoef**’ object. More specifically, the following attributes are relevant for function **spar()**:

- **nscreen** integer giving the number of variables to be retained after screening; defaults to $2n$.
- **split_data_prop**, double between 0 and 1 which indicates the proportion of the data that should be used for computing the screening coefficient. The remaining data will be used for estimating the marginal models in the SPAR algorithm; defaults to **NULL**. In this case the whole data will be used for estimating the screening coefficient and the marginal models.
- **type** character – either **"prob"** (indicating that probabilistic screening should be employed) or **"fixed"** (indicating that a fixed set of **nscreen** variables should be employed across the ensemble); defaults to **type = "prob"**.
- **reuse_in_rp** logical – indicates whether the screening coefficient should be reused at a later stage in the construction of the random projection. Defaults to **FALSE**.

The **control** argument, on the other hand, is a list containing extra parameters to be passed to the main function computing the screening coefficients.

The following screening coefficients are implemented in **spar**:

- **screen_marglik()** - computes the screening coefficients by the coefficient of x_j for $j = 1, \dots, p$ in a univariate GLM using the **stats::glm()** function.

$$\hat{\omega}_j =: \operatorname{argmin}_{\beta_j \in \mathbb{R}} \min_{\beta_0 \in \mathbb{R}} \sum_{i=1}^n -\ell(\beta_0, \beta_j; y_i, x_{ij})$$

It allows to pass a list of controls through the **control** argument to **stats::glm** such as weights, family, offsets.

- **screen_cor()** – computes the screening coefficients by the correlation between y and x_j using the function **stats::cor()**. It allows to pass a list of controls through the **control** argument to **stats::cor**.

- `screen_glmnet()` – computes by default the ridge coefficient where the penalty λ is very small (see [Parzer et al. 2024a](#), for clarification).

$$\hat{\omega} =: \operatorname{argmin}_{\beta \in \mathbb{R}^p} \min_{\beta_0 \in \mathbb{R}} \sum_{i=1}^n -\ell(\beta; y_i, x_i) + \frac{\varepsilon}{2} \sum_{j=1}^p \beta_j^2, \varepsilon > 0$$

The function relies on `glmnet::glmnet()` and, while it assumes by default $\alpha = 0$ and a small penalty, it allows to pass a list of controls through the `control` argument to `glmnet::glmnet()` such as `alpha = 1`. This screening coefficient is used as a default if `screencoef = NULL` in function call of `spar()` or `spar.cv()`.

All implemented `screen_*` functions return an object of class ‘`screencoef`’ which in turn is a list with three elements:

- a character `name`,
- `generate_fun()` – an R function for generating the screening coefficient. This function should have as following arguments: `x` – the matrix of standardized predictors – and `y` – the vector of (standardized in the Gaussian case) responses, and the argument `object`, which is a ‘`screencoef`’ object itself. It returns a vector of screening coefficients of length p .
- `control`, which is the control list passed by the user in `screen_*`. These controls are arguments which are needed in `generate_fun()` in order to generate the desired screening coefficients.

For illustration purposes, consider the object created by calling `screen_marglik()`:

```
R> obj <- screen_marglik()
```

A user-friendly `print` of the ‘`screencoef`’ is provided:

```
R> obj
#> Name: screen_marglik
#> Main attributes:
#> * proportion of data used for screening: 1
#> * number of screened variables: not provided, will default to 2n
#> * type: probabilistic screening
#> * screening coefficients: not yet computed from the data.
```

The structure of the object is the following:

```
R> unclass(obj)
#> $name
#> [1] "screen_marglik"
#>
#> $generate_fun
#> function(y, x, object) {
#>   if (is.null(object$control$family)) {
#>     object$control$family <- attr(object, "family")
#>   }
#>   coefs <- apply(x, 2, function(xj){
```

```

#>   glm_res <- do.call(function(...) glm(y ~ xj, ...),
#>                         object$control)
#>   glm_res$coefficients[2]
#> })
#>   coefs
#> }
#> <environment: namespace:spar>
#>
#> $control
#> list()
#>
#> attr("type")
#> [1] "prob"
#> attr("reuse_in_rp")
#> [1] FALSE

```

Function `generate_fun()` defines the generation of the screening coefficient. Note that it considers the controls in `object$control` when calling the `stats::glm()` function (unless it is provided, the `family` argument in `stats::glm()` will be set to the “global” family of the SPAR algorithm which is assigned inside the `spar()` function an attribute for the ‘`screencoef`’ object).

For convenience, a constructor function `constructor_screencoef()` is provided, which can be used to create new `screen_*` functions.

An example is presented in Section 4.1.

3.3. Random projections

Similar to the screening procedure, the objects for creating random projections are implemented as S3 classes ‘`randomprojection`’ and are created by functions `rp_*(..., control = list())`, which take `...` and a list of controls `control` as arguments.

Arguments related to the random projection can be passed through `...`, which will then be saved as attributes of the ‘`randomprojection`’ object. More specifically, the following attributes are relevant in the SPAR algorithm:

- **mslow**: integer giving the minimum dimension to which the predictors should be projected; defaults to $\log(p)$.
- **msup**: integer giving the maximum dimension to which the predictors should be projected; defaults to $n/2$.
- **data**: boolean indicating whether the random projection is data-driven.

Note that for random projection matrices which satisfy the JL lemma, **mslow** can be determined by employing existing results which give a lower bound on the goal dimension in order to preserve the distances between all pairs of points within a factor $(1 \pm \epsilon)$. For example, Achlioptas (2003) show $m_0 = \log n(4 + 2\tau)/(\epsilon^2/2 - \epsilon^3/3)$ for probability $1 - n^{-\tau}$.

The following random projections are implemented in **spar**:

- `rp_gaussian()` – random projection object where the generated matrix will have iid entries from a normal distribution (defaults to standard normal entries)
- `rp_sparse()` – random projection object where the generated matrix will be the one in (Achlioptas 2003) with `psi = 1` by default.
- `rp_cw()` – sparse embedding random projection in Clarkson and Woodruff (2013) for `rp_cw(data = FALSE)`. Defaults to `rp_cw(data=TRUE)`, which replaces the random elements on the diagonal by the ridge coefficients with a small penalty, as introduced in Parzer *et al.* (2024a).

The `rp_*` functions return an object of class ‘`randomprojection`’ which is a list with three elements:

- `name`,
- `generate_fun()` function for generating the random projection matrix. This function should have arguments `rp`, which is itself a ‘`randomprojection`’ object, `m`, the target dimension and a vector of indices `included_vector` which indicates the column index of the original variables in the `x` matrix to be projected using the random projection. This is needed due to the fact that screening can be employed pre-projection. It can return a matrix or a sparse matrix of class ‘`dgCMatrix`’ of the **Matrix** package (Bates, Maechler, and Jagan 2024) with `m` rows and `length(included_vector)` columns.
- `update_data_rp()` optional function used for data-driven random projections, which updates the ‘`randomprojection`’ object with data information which is relevant for the random projection. The updating happens only once, before the start of the SPAR algorithm, where appropriate attributes are added to the ‘`randomprojection`’ object. All relevant quantities are to be used in the `generate_fun()` function. This function should have arguments `rp`, which is a ‘`randomprojection`’ object to be updated, `x` – the matrix of standardized predictors, `y` – the vector of (standardized in the Gaussian case) responses and `screencoef` – a ‘`screencoef`’ object relevant for case where the screening coefficient should be reused in constructing the projection matrix. Returns a ‘`randomprojection`’ object.
- `update_rpm_w_data()` optional function for updating the random projection matrices provided in the argument `RPMs` of functions `spar()` and `spar.cv()` with data-dependent parameters. While `update_data_rp` is employed only once at the start of the algorithm, `update_rpm_w_data` specifies how to modify each random projection provided in `RPMs`. This is particularly relevant for the cross-validation procedure, which employs the random projection matrices generated by calling the `spar()` function on the whole data set before starting the cross-validation exercise. For example, in our implementation of the data-driven `rp_cw()`, in each fold, we only update the list of `RPMs` by adjusting the diagonal elements to the vector of screening coefficients computed on the current training data for the current fold, but do not modify the random elements in each fold, to reduce the computational burden. Defaults to `NULL`. If not provided, the values of the provided `RPMs` do not change.
- `control`, which is the control list in `rp_*`. These controls are arguments needed in `generate_fun()` in order to generate the desired random projection.

For illustration purposes, consider the implemented function `rp_gaussian()`, which generates a random projection with entries drawn from the normal distribution. The `print` method returns key information about the random projection procedure.

```
R> obj <- rp_gaussian()
R> obj
#> Name: rp_gaussian
#> Main attributes:
#> * Data-dependent: FALSE
#> * Lower bound on goal dimension m: not provided, will default to log(p).
#> * Upper bound on goal dimension m: not provided, will default to n/2.
```

We turn to looking at the structure of the object:

```
R> unclass(obj)
#> $name
#> [1] "rp_gaussian"
#>
#> $generate_fun
#> function(rp, m, included_vector) {
#>   p <- length(included_vector)
#>   control_rnorm <-
#>     rp$control[names(rp$control) %in% names(formals(rnorm))]
#>   vals <- do.call(function(...)
#>     rnorm(m * p, ...), control_rnorm)
#>   RM <- matrix(vals, nrow = m, ncol = p)
#>   return(RM)
#> }
#> <environment: namespace:spar>
#>
#> $update_data_fun
#> NULL
#>
#> $update_rpm_w_data
#> NULL
#>
#> $control
#> list()
#>
#> attr("data")
#> [1] FALSE
```

The `generate_fun()` function returns a matrix with `m` rows and `length(included_vector)` columns. Note that `included_vector` gives the indices of the variables which have been selected by the screening procedure. In this case, where the random projection does not use any data information, we are only interested in the length of this vector.

The functions `update_data_fun()` and `update_rpm_w_data()` are `NULL` as this conventional random projection is data-agnostic.

3.4. Marginal models

The package provides a class ‘`sparmodel`’ for the marginal model to be fitted for each element of the ensemble. The framework currently assumes that the linear predictor is a linear combination of the projected variables.

Similar to the objects for random projection and screening coefficients, the functions which create these objects have arguments `...` (to be saved as attributes) and `control` (to be used in the main function for building the model).

The two functions implemented are `spar_glmnet()`, which allows regularized GLMs as marginal models using function `glmnet::glmnet()` (where the default is to estimate a ridge regression with the small penalty value), and `spar_glm()` which estimates unregularized GLMs using `stats::glm.fit()`.

An object of class ‘`sparmodel`’ is a list with elements:

- `name`,
- `model_fun()` – a function which takes `y` (the vector of standardized responses), `z` (the matrix of reduced predictors) and a further argument which is the object of class ‘`sparmodel`’ itself.
- `update_model()` – an optional function which can add further attributes to the ‘`sparmodel`’ object which is called at the beginning of the SPAR algorithm. In the case of `spar_glmnet()` this function manipulates the ‘`family`’ object in a way which is convenient for `glmnet::glmnet()`.¹

The default is to use `spar_glm()` for Gaussian family with identity link and `spar_glmnet()` for the other families.

3.5. Methods

Methods `print`, `plot`, `coef`, `predict` are available for both ‘`spar`’ and ‘`spar.cv`’ classes.

print

The `print` method returns information on ν_{best} M_{best} , the number of active predictors (i.e., predictors which have at least a nonzero coefficient across the marginal models) and a five-point summary of the non-zero coefficients.

```
R> set.seed(12)
R> spar_res <- spar(example_data$x, example_data$y,
+                 xval = example_data$xtest,
+                 yval = example_data$ytest,
+                 nummods = c(5, 10, 15, 20, 25, 30))
R> spar_cv <- spar.cv(example_data$x, example_data$y,
```

¹In the case of families Gaussian, binomial and Poisson with canonical link, the family object is replaced by a string containing the name of the family. This leads to `glmnet` using the faster specialized algorithms rather than the general algorithm implemented for all `[family].class` objects.


```

+               nummods = c(5, 10, 15, 20, 25, 30))

R> spar_res
#> spar object:
#> Smallest Validation Measure reached for nummod=30,
#>               nu=1.72e-02 leading to 1002 / 2000 active predictors.
#> Summary of those non-zero coefficients:
#>      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
#> -1.19688 -0.07144  0.02134  0.02814  0.12008  0.99405

```

For ‘spar.cv’ it also provides the same information for the (ν, M) combination chosen by the one-standard error rule.

```

R> spar_cv
#> spar.cv object:
#> Smallest CV-Meas 2230.9 reached for nummod=30,
#>               nu=0.00e+00 leading to 1825 / 2000 active predictors.
#> Summary of those non-zero coefficients:
#>      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
#> -0.825363 -0.041159  0.004422  0.018693  0.072914  0.953501
#>
#> Sparsest coefficient within one standard error of best CV-Meas
#>               reached for nummod=5, nu=5.40e-03
#> leading to 1003 / 2000 active
#>               predictors with CV-Meas 2748.4.
#> Summary of those non-zero coefficients:
#>      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
#> -1.01802 -0.10460  0.04862  0.03436  0.17091  1.32992

```

coef

Method `coef` takes as inputs a ‘spar’ or ‘spar.cv’ object, together with further arguments:

- **nummod** – number of models used to compute the averaged coefficients; value of **nummod** with minimal **measure** is used if not provided.
- **nu** – threshold level used to compute the averaged coefficients; value with minimal **measure** is used if not provided.

```

R> str(coef(spar_res))
#> List of 4
#> $ intercept: num 2.86
#> $ beta      : num [1:2000] 0 0 0.746 0 0 ...
#> $ nummod    : num 30
#> $ nu        : num 0.0172

```

It returns a list with the intercept, vector of **beta** coefficients and the **nummod** and **nu** employed in the calculation.

Additionally for ‘spar.cv’, the `coef` method also has argument **opt_par** which is one of

`c("lse","best")` and chooses whether to select the best pair of `nu`s and `nummods` according to cross-validation `measure`, or the solution yielding sparsest vector of coefficients within one standard deviation of that optimal cross-validation `measure`. This argument is ignored when `nummod` and `nu` are given.

predict

Functionality for computing predictions is provided through the method `predict` which takes a `'spar'` or `'spar.cv'` object, together with

- `xnew` – matrix of new predictor variables; must have same number of columns as `x`.
- `type` – the type of required predictions; either on `"response"` level (default) or on `"link"` level
- `avg_type` – type of averaging used across the marginal models; either on `"link"` (default) or on `"response"` level
- `nummod` – number of models used to compute the averaged coefficients; value of `nummod` with minimal `measure` is used if not provided.
- `nu` – threshold level used to compute the averaged coefficients; value with minimal `measure` is used if not provided.
- `coef` – optional vector of coefficients to be used directly in the prediction.

Additionally, for class `'spar.cv'`, argument `opt_par` is available and used in the computation of the coefficients to be used for prediction (see above description of method `coef`).

plot

Plotting functionality is provided through the `plot` method, which takes a `'spar'` or `'spar.cv'` object, together with further arguments:

- `plot_type` – one of:
- `"Val_Measure"` plots the (cross-)validation `measure` for either a grid of `nu` values for a fixed number of models `nummod` or viceversa.
- `"Val_numAct"` plots the number of active variables for either a grid of `nu` values for a fixed number of models `nummod` or viceversa.
- `"res-vs-fitted"` produces a residuals-vs-fitted plot. The residuals are computed as $y - \hat{y}$, where $\hat{y}$ is the prediction computed on response level.
- `"coefs"` produces a plot of the value of the standardized coefficients for each predictor in each marginal model (before thresholding). For each predictor, the values of the coefficients are sorted from largest to smallest across the marginal models and then represented in the plot.
- `plot_along` – one of `c("nu","nummod")`; for `plot_type = "Val_Measure"` or `plot_type = "Val_numAct"` it indicates whether the values of the cross-validation measure or number of active variables, respectively, should be shown for a grid of ν

values while keeping the number of models `nummod` fixed or viceversa. This argument is ignored when `plot_type = "res-vs-fitted"` or `plot_type = "coefs"`.

- `nummod` – fixed value for number of models when `plot_along = "nu"` for `plot_type = "Val_Measure"` or `"Val_numAct"`; if `plot_type = "res-vs-fitted"`, it is used in the `predict` method, as described above.
- `nu` – fixed value for ν when `plot_along = "nummod"` for `plot_type = "Val_Measure"` or `"Val_numAct"`; if `plot_type = "res-vs-fitted"`, it is used in the `predict` method, as described above.
- `xfit` – if `plot_type = "res-vs-fitted"`, it is the matrix of predictors used in computing the fitted values. This argument must be provided for the plot of residuals and fitted values, as the `'spar'` or `'spar.cv'` objects do not store the original data.
- `yfit` – if `plot_type = "res-vs-fitted"`, vector of responses used in computing the residuals. This argument must be provided for the plot of residuals and fitted values, as the `'spar'` or `'spar.cv'` objects do not store the original data.
- `prange` – optional vector of length 2 in case `plot_type = "coefs"` which gives the limits of the predictors' plot range; defaults to `c(1, p)`.
- `coef_order` – optional index vector of length p in case `plot_type = "coefs"` to give the order of the predictors; defaults to `1 : p`.

For class `'spar.cv'` there is the extra argument `opt_par = c("best", "lse")` which, for `plot_type = "res-vs-fitted"` indicates whether the predictions should be based on coefficients using the best (ν, M) combination or on the combination which delivers the sparsest β having validation measure within one standard deviation from the minimum

The `plot` methods return objects of class `'ggplot'` (Wickham 2016).

4. Extensibility

4.1. Screening coefficients

We exemplify how new screening coefficients implemented in package **VariableScreening** can easily be used in the framework of **spar**.

We start by defining the function for generating the screening coefficients using the `screenIID()` function in **VariableScreening**.

```
R> generate_scr_sirs <- function(y, x, object) {
+   res_screen <- do.call(function(...)
+     VariableScreening::screenIID(x, y, ...),
+     object$control)
+   coefs <- res_screen$measurement
+   coefs
+ }
```

Note that `screenIID()` also takes `method` as an argument. To allow for flexibility, we do not fix the method in `generate_scr_sirs()` but rather allow the user to pass a method through the `control` argument in the `screen_*` function. This function is created using the helper `constructor_screencoef()`:

```
R> screen_sirs <- constructor_screencoef(
+   "screen_sirs",
+   generate_fun = generate_scr_sirs)
```

We now call the `spar()` function with the newly created screening procedure. We consider the method SIRS of [Zhu et al. \(2011\)](#), which ranks the predictors by their correlation with the rank-ordered response and we do not perform probabilistic variable screening but employ the top $2n$ variables in each marginal model.

```
R> set.seed(123)
R> spar_example <- spar(example_data$x, example_data$y,
+   screencoef = screen_sirs(type = "fixed",
+   control=list(method = "SIRS")),
+   measure = "mse")
R> spar_example
#> spar object:
#> Smallest Validation Measure reached for nummod=20,
#>          nu=1.75e-03 leading to 396 / 2000 active predictors.
#> Summary of those non-zero coefficients:
#>      Min.    1st Qu.    Median      Mean    3rd Qu.      Max.
#> -0.6918230 -0.0928444  0.0009116  0.0943062  0.1777889  2.0089709
```

4.2. Random projections

We exemplify how new random projections can be implemented in the framework of **spar**.

We implement the random projection of [Cannings and Samworth \(2017\)](#), who propose using the Haar measure for generating the random projections. They simulate matrices from the Haar measure by independently drawing each entry of a matrix Q from a standard normal distribution, and then to take the projection matrix to be the transpose of the matrix of left singular vectors in the singular value decomposition of Q . Moreover, they suggest using “good” random projections, in the sense that they deliver the best out-of-sample prediction. The proposed approach employs B_1 models in an ensemble of classifiers and for each model k , B_2 data independent random projections are generated and the one with the lowest error on a test set is the one chosen to project the variables in model k .

We can implement such a random projection in **spar** by the following building block:

```
R> update_data_cannings <- function(rp, x, y, screencoef = NULL) {
+   attr(rp, "dataset") <- list(x = x, y = y)
+   rp
+ }
```

This is the function which adds data information to the random projection object. Here, the whole data can be added as information for the M random projection (alternatively, one could

only pass sufficient statistics for computing the desired measures).

While the B_2 random projections are data-agnostic, the `generate_fun()` element of the random projection will need the data information in order to evaluate which method performs best in terms of an error measure. We will, in the following, define the function for the generation of the random projection matrix to be used in a model k .

This helper simulates $m \times p$ matrices from the Haar measure:

```
R> simulate_haar <- function(m, p) {
+   R0 <- matrix(1/sqrt(p) * rnorm(p * m), nrow = p, ncol = m)
+   RM <- qr.Q(qr(R0))[, seq_len(m)]
+   t(RM)
+ }
```

The function that generates the random projection matrix for model k uses 25% of the data as a test set. After estimating a ridge regression with a minimal penalty on the training data (this is the marginal model to be employed in the SPAR algorithm), it chooses the best among B_2 random projections in terms of minimizing misclassification error for the binomial family and MSE for all other families on the test set:

```
R> generate_cannings <- function(rp, m, included_vector) {
+   x <- attr(rp, "dataset")$x[, included_vector]
+   y <- attr(rp, "dataset")$y
+   n <- nrow(x); p <- ncol(x)
+   B2 <- ifelse(is.null(rp$control$B2), 50, rp$control$B2)
+   id_test <- sample(n, size = n %/% 4)
+   xtrain <- x[-id_test, ]; xtest <- x[id_test,]
+   ytrain <- y[-id_test]; ytest <- y[id_test]
+   if (is.null(rp$control$family)) rp$control$family <- attr(rp, "family")
+   if (is.null(rp$control$alpha)) rp$control$alpha <- 1
+   control_glmnet <-
+     rp$control[names(rp$control) %in% names(formals(glmnet::glmnet))]
+   best_val <- 1e5
+   for (b in seq_len(B2)) {
+     RM <- simulate_haar(m, p)
+     xrp <- tcrossprod(xtrain, RM)
+     mod <- do.call(function(...)
+       glmnet::glmnet(x = xrp, y = ytrain, ...), control_glmnet)
+     coefs <- coef(mod, s = min(mod$lambda))
+     eta_test <- (cbind(1, tcrossprod(xtest, RM)) %*% coefs)
+     pred <- rp$control$family$linkinv(as.vector(eta_test))
+     out_perf <- ifelse(rp$control$family$family == "binomial",
+       mean(((pred > 0.5) + 0) != ytest),
+       mean((pred - ytest)^2))
+     if (out_perf < best_val) {
+       best_val <- out_perf
+       best_RM <- RM
+     }
+   }
```

```

+   rm(RM)
+ }
+ # error_all <- lapply(seq_len(B2), FUN = function(s){
+ #   RM <- simulate_haar(m, p)
+ #   xrp <- tcrossprod(xtrain, RM)
+ #   mod <- do.call(function(...)
+ #     glmnet::glmnet(x = xrp, y = ytrain, ...), control_glmnet)
+ #   coefs <- coef(mod, s = min(mod$lambda))
+ #   eta_test <- (cbind(1, tcrossprod(xtest, RM)) %*% coefs)
+ #   pred <- rp$control$family$linkinv(as.vector(eta_test))
+ #   out <- ifelse(rp$control$family$family == "binomial",
+ #     mean(((pred > 0.5) + 0) != ytest),
+ #     mean((pred - ytest)^2))
+ #   list(RM, out)
+ # })
+ # id_best <- which.min(sapply(error_all, "[", 2))
+ # RM <- error_all[[id_best]][[1]]
+ return(best_RM)
+}

```

In the cross-validation procedure, we do not generate new matrices for each step to keep computational costs low, so we do not specify a function `update_rpm_w_data()`.

Putting it all together, we get:

```

R> rp_cannings <- constructor_randomprojection("rp_cannings",
+   generate_fun = generate_cannings,
+   update_data_fun = update_data_cannings
+)

```

We can now estimate SPAR for a binomial model, where we transform the response to a binary variable.

```

R> ystar <- (example_data$y > 0) + 0
R> ystarval <- (example_data$ytest > 0) + 0

```

We use 50 models (which is in line to recommendations in [Cannings and Samworth \(2017\)](#)), and no screening procedure. Note that, if screening should not be performed, `nscreen` can be set to `p` in the `screen_*` function. Moreover, we do not perform any thresholding.

```

R> set.seed(12345)
R> spar_example_1 <- spar(x = example_data$x, y = ystar,
+   family = binomial(),
+   screencoef = screen_marglik(nscreen = ncol(example_data$x)),
+   rp = rp_cannings(control = list(B2 = 50, lambda.min.ratio = 0.01)),
+   nus = 0, nummods = 20,
+   xval = example_data$xtest, yval = ystarval,
+   measure = "class"
+)

```

Using the data-driven `rp_cw()`:

```

R> set.seed(12345)
R> spar_example_2 <- spar(x = example_data$x, y = ystar,
+                         family = binomial(),
+                         screencoef = screen_marglik(nscreen = ncol(example_data$x)),
+                         rp = rp_cw(data = TRUE),
+                         nus = 0, nummods = 50,
+                         xval = example_data$xtest, yval = ystarval,
+                         measure = "class"
+ )

```

We can now compare the two approaches by looking at the minimum measure `Meas` achieved on the validation set:

```

R> spar_example_1$val_res[which.min(spar_example_1$val_res$Meas),]
#>   nnu nu nummod numAct Meas
#> 1   1  1  0      50   2000 0.17
R> spar_example_2$val_res[which.min(spar_example_2$val_res$Meas),]
#>   nnu nu nummod numAct Meas
#> 1   1  1  0      50   2000 0.17

```

5. Illustrations

5.1. Face image data

We illustrate the functionality of **spar** on the Isomap data set containing $n = 698$ black and white face images of size $p = 64 \times 64 = 4096$ together with the faces' horizontal looking direction angle as the response variable. The data set is publicly available for download:

```

R> url1 <- "https://web.archive.org/web/20150922051706/"
R> url2 <- "http://isomap.stanford.edu/face_data.mat.Z"
R> download.file(paste0(url1, url2),
+               file.path("face_data.mat.Z"))
R> system('uncompress face_data.mat.Z')

```

The `.mat` file format can be read using **R.matlab** ([Bengtsson 2022](#)):

```

R> library("R.matlab")
R> facedata <- readMat(file.path("face_data.mat"))
R> x <- t(facedata$images)
R> y <- facedata$poses[1,]

```

We can visualize one observation in this data set in the left panel Figure 3 (code for reproducing the figure is provided in the supplementary materials).

This data set has the issue of many columns being almost constant. Functions `spar()` and `spar.cv()` ignore constant columns from the estimation, but we can further alleviate the issue by setting all columns which have a low standard deviation to zero.

```

R> x[, apply(x, 2, sd) < 0.01] <- 0

```

We split the data into training vs test sample

```
R> set.seed(123)
R> ntot <- length(y); ntest <- ntot * 0.25
R> testind <- sample(ntot, ntest, replace = FALSE)
R> xtrain <- as.matrix(x[-testind, ]); ytrain <- y[-testind]
R> xtest <- as.matrix(x[testind, ]); ytest <- y[testind]
```

Versions of the SPAR algorithm have been shown to achieve a high prediction performance on this data set (see [Parzer et al. 2024b](#)). We illustrate here the SPAR algorithm with cross-validation, where the data-driven random projection in [Parzer et al. \(2024b\)](#) is used together with screening based on the ridge coefficients with the penalty approaching zero ([Wang and Leng 2016](#)). To ensure convergence of the **glmnet** algorithm in computing the screening coefficient (which is also employed in the data-driven random projection), we set the `lambda.min.ratio` parameter to 0.001. Moreover, each marginal model is a linear regression with the regression coefficients estimated by maximum likelihood.

```
R> library("spar")
R> set.seed(123)
R> control_glmnet <- list(lambda.min.ratio = 0.001)
R> spar_faces <- spar.cv(xtrain, ytrain,
+   model = spar_glm(),
+   screencoef = screen_glmnet(control = control_glmnet, reuse_in_rp = TRUE),
+   rp = rp_cw(data = TRUE),
+   nummods = c(5, 10, 20, 50), measure = "mse")
```

The print method returns:

```
R> spar_faces
#> spar.cv object:
#> Smallest CV-Meas 12.6 reached for nummod=50,
#>          nu=0.00e+00 leading to 3689 / 4096 active predictors.
#> Summary of those non-zero coefficients:
#>      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
#> -652.8332  -0.0822    0.0000    0.6897    0.0945 1291.5781
#>
#> Sparsest coefficient within one standard error of best CV-Meas
#>          reached for nummod=5, nu=3.45e-03
#> leading to 1597 / 4096 active
#>          predictors with CV-Meas 17.1.
#> Summary of those non-zero coefficients:
#>      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
#> -826.8860  -0.3638    0.0200    0.5223    0.4518  966.2025
```

The `plot` method for ‘`spar.cv`’ objects displays by default the measure employed in the cross-validation (in this case MSE) for a grid of thresholding values ν , where the number of models is fixed to the value found to perform best in cross-validation exercise (see [Figure 1](#)).

```
R> plot(spar_faces)
```

We observe that 50 models and $\nu = 0$ delivers the lowest MSE but that the differences in




Figure 1: Plot of mean squared error over a grid of threshold values ν for fixed number of optimal models $M = 20$ for the *Isomap faces* data set. The red points correspond to the threshold which achieves the lowest cross-validation measure and the one with the largest cross-validation measure within one standard deviation away from minimum. Confidence bands represent one standard deviation in the measures across the number of folds.

MSE compared to the value of $\nu = 0.00408$ chosen by 1-standard-error rule (for 50 models) are minimal.

The coefficients of the different variables (in this example pixels) obtained by averaging over the coefficients the marginal models (for optimal ν and number of models) are given by:

```
R> face_coef <- coef(spar_faces, opt_par = "best")
R> str(face_coef)
#> List of 4
#> $ intercept: num -24.2
#> $ beta      : num [1:4096] 0.5945 0.1556 0.0303 0 -0.7395 ...
#> $ nummod    : num 50
#> $ nu        : num 0
```

For a sparser solution we can compute the coefficients using `opt_par = "1se"` which leads to more sparsity and a lower number of models.

```
R> face_coef_1se <- coef(spar_faces, opt_par = "1se")
R> str(face_coef_1se)
#> List of 4
#> $ intercept: num -25.5
#> $ beta      : num [1:4096] -6.66 0 0 0 0 ...
#> $ nummod    : num 5
#> $ nu        : num 0.00345
```

The standardized coefficients from each of `max(nummods)` models (before thresholding) can be plotted by setting `plot_type = "coefs"` (see Figure 2).

```
R> plot(spar_faces, plot_type = "coefs")
```

We observe that pixels close to each other have more correlation than pixels further apart.

The `predict()` function can be applied to the `'spar.cv'` object. We will employ the sparser solution chosen by the `opt_par = "1se"` rule:




Figure 2: Coefficient plot for all variables and all $M = 50$ models in the SPAR ensemble for the *Isomap faces* data set. The coefficients are standardized, before thresholding.

```
R> ynew <- predict(spar_faces, xnew = xtest, coef = face_coef_1se)
```

We can easily compare the predictive performance of the estimated version of the SPAR algorithm with another version proposed in [Mukhopadhyay and Dunson \(2020\)](#), where the sparse random projection in [Achlioptas \(2003\)](#) is combined with probabilistic marginal correlation screening.

```
R> tarp_faces <- spar.cv(xtrain, ytrain,
+   model = spar_glm(),
+   screencoef = screen_cor(),
+   rp = rp_sparse(psi = 1/3),
+   nummods = c(5, 10, 20, 50), measure = "mse")
R> ynew_tarp <- predict(tarp_faces, xnew = xtest,
+   coef = coef(tarp_faces, opt_par = "1se"))
R> c("SPAR" = mean((ytest-ynew)^2),
+   "TARP" = mean((ytest-ynew_tarp)^2))
#>      SPAR      TARP
#> 39.39278 67.81535
```

For this data set, one can visualize $\hat{\beta}_j^{1se} x_{i,j}^{new}$, the effect of each pixel in predicting the face orientation in a given image. The contribution of each pixel in predicting the orientation of the face in the left panel of Figure 3 can be visualized in the right panel of Figure 3. We observe that the face contour and the hairline are most informative for the prediction.

5.2. Darwin data set

The Darwin data set ([Cilia, De Gregorio, De Stefano, Fontanella, Marcelli, and Parziale 2022](#)) contains a binary response for Alzheimer's disease (AD) together with extracted features from 25 handwriting tests (18 features per task) for 89 AD patients and 85 healthy people




Figure 3: Left: Image corresponding to one observation in the *Isomap faces* data set. Right: The effect of each pixel $\hat{\beta}_j^{\text{lse}} x_{i,j}^{\text{new}}$ in predicting the face orientation in left panel.

$(n = 174)$.²

```
R> download.file("https://archive.ics.uci.edu/static/public/732/darwin.zip",
+               "darwin.zip")
```

```
R> darwin_tmp <- read.csv(unzip("darwin.zip", "data.csv"),
+                        stringsAsFactors = TRUE)
```

Before proceeding with the analysis, the data is screened for multivariate outliers using the DDC algorithm in package **cellWise** (Raymaekers and Rousseeuw 2023).

```
R> darwin_orig <- list(
+   x = darwin_tmp[, !(colnames(darwin_tmp) %in% c("ID", "class"))],
+   y = as.numeric(darwin_tmp$class) - 1)
R> tmp <- cellWise::DDC(
+   as.matrix(darwin_orig$x),
+   list(returnBigXimp = TRUE,
+        tolProb = 0.999,
+        silent = TRUE))
#>
#> The final data set we will analyze has 174 rows and 446 columns.
#>
R> darwin <- list(x = tmp$Ximp, y = darwin_orig$y)
```

The structure of the data is:

```
R> str(darwin)
#> List of 2
#> $ x: num [1:174, 1:450] 5160 3721 2600 2130 2310 ...
#> ..- attr(*, "dimnames")=List of 2
#> .. ..$ : NULL
```

²The data set can be downloaded from <https://archive.ics.uci.edu/dataset/732/darwin>




Figure 4: Average number of active variables for the grid of thresholding values ν and $M = 10$ models for the *Darwin* data set. The red points correspond to the average number of active variables for the model with the lowest cross-validation measures and to the one chosen by the 1-standard-error rule.

```
#> .. ..$ : chr [1:450] "air_time1" "disp_index1" "gmrt_in_air1" "gmrt_on_paper1" ...
#> $ y: num [1:174] 1 1 1 1 1 1 1 1 1 1 ...
```

We estimate the SPAR algorithm with the screening and random projection introduced in [Parzer et al. \(2024a\)](#) for binomial family and logit link, using 1–area under the ROC curve as the cross-validation measure:

```
R> spar_darwin <- spar.cv(darwin$x, darwin$y,
+                          family = binomial(logit),
+                          nummods = c(5, 10, 20, 50),
+                          measure = "1-auc")
```

We can look at the average number of active variables for a grid of ν where the number of models is fixed to the value found to perform best in cross-validation exercise (see Figure 4).

```
R> plot(spar_darwin, plot_type = "Val_numAct")
```

We observe again that no thresholding achieves the best measure and translates to almost all variables being active (some variables can be inactive at $\nu = 0$ as they may never be screened). The 1-standard-error rule would however indicate that more sparsity can be introduced without too much increase in the cross-validation measure.

Finally, coefficients of the predictors across the maximum number of considered marginal models (in this case $M = 50$) can be visualized with `plot(spar_darwin, plot_type = "coefs")`. In this data set, the predictors are ordered by task, where the first 18 covariates represent different features measured for the first task. Given that there is clear grouping in the variables in this example, we can reorder the coefficients for plotting by grouping them by feature, rather than task. This allows to assess how the different features (e.g., time it takes to complete a certain task) relate to the likelihood of having AD and how stable the sign and magnitude of the coefficient is across the models in the ensemble. We can achieve




Figure 5: Coefficient plot for all variables and all considered $M = 50$ models in the SPAR ensemble *Darwin* data set. The coefficients are standardized, before thresholding.

this by using reordering argument `coef_order` in method `plot` with `plot_type = "coefs"` (see Figure 5 which can be reproduced with code in the supplementary materials).

In general we observe that the different features measures across different tasks have the same impact on the probability of AD (observable by the blocks of blue or red lines).

6. Conclusion

Package **spar** can be employed for modeling data in a high-dimensional setting, where the number of predictors is much higher than the number of observations. The package provides an implementation an algorithm for sparse projected and average regression (SPAR) proposed in Parzer *et al.* (2024a) which combines variable screening and random projection in an ensemble of GLMs. The package provides flexible classes for i) specifying the coefficient based on which screening should be performed (both in a classical fashion, where the predictors with the highest screening coefficient are selected for subsequent or in a probabilistic fashion, where variables are sampled for inclusion with probabilities proportional to their screening coefficient), ii) generating the random projection to be employed in each marginal model. Screening coefficients based on marginal correlation between the predictors and the response, marginal coefficients from a GLM or ridge coefficients are provided in the package. Moreover, several random projections are implemented: the Gaussian and sparse matrices which are data-agnostic and satisfy the JL lemma and the data-driven projection proposed in Parzer *et al.* (2024b) for linear regression and extended to GLMs in Parzer *et al.* (2024a). This method has the advantage of approximately capturing the true β in the span of the random projection matrix, i.e., it ensures that the true regression coefficients can be recovered approximately after the projection. Methodologically, the SPAR algorithm, particularly when paired with the data-driven random projection in Parzer *et al.* (2024a), has been demonstrated to perform effectively across different degrees of sparsity of the coefficient vector and to offer competitive predictions and variable ranking in both sparse and dense settings.

The flexibility and adaptability of the **spar** package make it an attractive choice for practi-

tioners and researchers. It encourages exploration of new methods for variable screening and random projections or the combination of existing approaches to tailor solutions to specific data requirements.

Computational details

The results in this paper were obtained using R 4.4.0.

R itself and all packages used are available from CRAN at <https://CRAN.R-project.org/>.

Acknowledgments

Roman Parzer and Laura Vana-Gür acknowledge funding from the Austrian Science Fund (FWF) for the project “High-dimensional statistical learning: New methods to advance economic and sustainability policies” (ZK 35), jointly carried out by WU Vienna University of Economics and Business, Paris Lodron University Salzburg, TU Wien, and the Austrian Institute of Economic Research (WIFO).

References

- Achlioptas D (2003). “Database-Friendly Random Projections: Johnson-Lindenstrauss with Binary Coins.” *Journal of Computer and System Sciences*, **66**(4), 671–687. ISSN 0022-0000. doi:[10.1016/S0022-0000\(03\)00025-4](https://doi.org/10.1016/S0022-0000(03)00025-4). Special Issue on PODS 2001.
- Aghila G, Siddharth R (2020). **RandPro**: *Random Projection with Classification*. doi:[10.32614/CRAN.package.RandPro](https://doi.org/10.32614/CRAN.package.RandPro). R package version 0.2.2.
- Ailon N, Chazelle B (2009). “The Fast Johnson–Lindenstrauss Transform and Approximate Nearest Neighbors.” *SIAM Journal on computing*, **39**(1), 302–322. doi:[10.1137/060673096](https://doi.org/10.1137/060673096).
- Bates D, Maechler M, Jagan M (2024). **Matrix**: *Sparse and Dense Matrix Classes and Methods*. doi:[10.32614/CRAN.package.Matrix](https://doi.org/10.32614/CRAN.package.Matrix). R package version 1.7-0.
- Bengtsson H (2022). **R.matlab**: *Read and Write MAT Files and Call MATLAB from Within R*. doi:[10.32614/CRAN.package.R.matlab](https://doi.org/10.32614/CRAN.package.R.matlab). R package version 3.7.0.
- Cannings TI, Samworth RJ (2017). “Random-projection ensemble classification.” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, **79**(4), 959–1035. doi:[10.1111/rssb.12228](https://doi.org/10.1111/rssb.12228).
- Cannings TI, Samworth RJ (2021). **RPEnsemble**: *Random Projection Ensemble Classification*. R package version 0.5, URL [10.32614/CRAN.package.RPEnsemble](https://CRAN.R-project.org/package=RPEnsemble).
- Cantoni E, Ronchetti E (2001). “Robust inference for generalized linear models.” *Journal of the American Statistical Association*, **96**(455), 1022–1030. doi:[10.1198/016214501753209004](https://doi.org/10.1198/016214501753209004).

- Cheng X, Wang H, Zhu L, Zhong W, Zhou H (2024). **MFSIS: Model-Free Sure Independent Screening Procedures**. doi:10.32614/CRAN.package.MFSIS. R package version 0.2.1.
- Cho H, Fryzlewicz P (2012). “High dimensional variable selection via tilting.” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, **74**(3), 593–622. doi:10.1111/j.1467-9868.2011.01023.x.
- Cilia ND, De Gregorio G, De Stefano C, Fontanella F, Marcelli A, Parziale A (2022). “Diagnosing Alzheimer’s disease from on-line handwriting: A novel dataset and performance benchmarking.” *Engineering Applications of Artificial Intelligence*, **111**, 104822. ISSN 0952-1976. doi:10.1016/j.engappai.2022.104822.
- Clarkson KL, Woodruff DP (2013). “Low Rank Approximation and Regression in Input Sparsity Time.” In *Proceedings of the Forty-Fifth Annual ACM Symposium on Theory of Computing*, STOC ’13, p. 81–90. Association for Computing Machinery, New York, NY, USA. ISBN 9781450320290. doi:10.1145/2488608.2488620.
- Cui H, Li R, Zhong W (2015). “Model-free feature screening for ultrahigh dimensional discriminant analysis.” *Journal of the American Statistical Association*, **110**(510), 630–641. doi:10.1080/01621459.2014.920256.
- Fan J, Feng Y, Song R (2011). “Nonparametric independence screening in sparse ultra-high-dimensional additive models.” *Journal of the American Statistical Association*, **106**(494), 544–557. doi:10.1198/jasa.2011.tm09779.
- Fan J, Feng Y, Wu Y (2010). “High-dimensional variable selection for Cox’s proportional hazards model.” In *Borrowing strength: Theory powering applications—a Festschrift for Lawrence D. Brown*, volume 6, pp. 70–87. Institute of Mathematical Statistics. doi:10.1214/10-IMSCOLL606.
- Fan J, Lv J (2007). “Sure Independence Screening for Ultra-High Dimensional Feature Space.” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, **B 70**. doi:10.1111/j.1467-9868.2008.00674.x.
- Fan J, Samworth R, Wu Y (2009). “Ultrahigh dimensional feature selection: beyond the linear model.” *The Journal of Machine Learning Research*, **10**, 2013–2038. URL <https://jmlr.csail.mit.edu/papers/v10/fan09a.html>.
- Fan J, Song R (2010). “Sure independence screening in generalized linear models with NP-dimensionality.” *The Annals of Statistics*, **38**(6), 3567 – 3604. doi:10.1214/10-AOS798.
- Frankl P, Maehara H (1988). “The Johnson-Lindenstrauss Lemma and the Sphericity of Some Graphs.” *Journal of Combinatorial Theory, Series B*, **44**(3), 355–362. ISSN 0095-8956. doi:10.1016/0095-8956(88)90043-3.
- Gataric M, Wang T, Samworth RJ (2019). **SPCAvRP: Sparse Principal Component Analysis via Random Projections (SPCAvRP)**. R package version 0.4, URL [10.32614/CRAN.package.SPCAvRP](https://CRAN.R-project.org/package=SPCAvRP).
- Johnson W, Lindenstrauss J (1984). “Extensions of Lipschitz Maps into a Hilbert Space.” *Contemporary Mathematics*, **26**, 189–206. doi:10.1090/conm/026/737400.

- Ke C (2023). “Sufficient Variable Screening With High-Dimensional Controls.” *Electronic Journal of Statistics*, **17**(2), 2139 – 2179. doi:[10.1214/23-EJS2150](https://doi.org/10.1214/23-EJS2150).
- Li P, Hastie TJ, Church KW (2006). “Very Sparse Random Projections.” In *Proceedings of the 12th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD ’06, p. 287–296. Association for Computing Machinery, New York, NY, USA. ISBN 1595933395. doi:[10.1145/1150402.1150436](https://doi.org/10.1145/1150402.1150436).
- Li R, Huang L, Dziak J (2022). **VariableScreening**: *High-Dimensional Screening for Semi-parametric Longitudinal Regression*. doi:[10.32614/CRAN.package.VariableScreening](https://doi.org/10.32614/CRAN.package.VariableScreening). R package version 0.2.1.
- Li R, Zhong W, Zhu L (2012). “Feature Screening via Distance Correlation Learning.” *Journal of the American Statistical Association*, **107**(499), 1129–1139. doi:[10.1080/01621459.2012.695654](https://doi.org/10.1080/01621459.2012.695654).
- Ma X, Zhang J (2016). “Robust Model-Free Feature Screening via Quantile Correlation.” *Journal of Multivariate Analysis*, **143**, 472–480. doi:[10.1016/j.jmva.2015.10.010](https://doi.org/10.1016/j.jmva.2015.10.010).
- Mai Q, Zou H (2013). “The Kolmogorov Filter for Variable Screening in High-Dimensional Binary Classification.” *Biometrika*, **100**(1), 229–234. doi:[10.1093/biomet/ass062](https://doi.org/10.1093/biomet/ass062).
- Mai Q, Zou H (2015). “The Fused Kolmogorov Filter: A Nonparametric Model-Free Screening Method.” *The Annals of Statistics*, **43**(4), 1471 – 1497. doi:[10.1214/14-AOS1303](https://doi.org/10.1214/14-AOS1303).
- Matoušek J (2008). “On Variants of the Johnson–Lindenstrauss Lemma.” *Random Structures & Algorithms*, **33**(2), 142–156. doi:[10.1002/rsa.20218](https://doi.org/10.1002/rsa.20218).
- Mukhopadhyay M, Dunson DB (2020). “Targeted Random Projection for Prediction From High-Dimensional Features.” *Journal of the American Statistical Association*, **115**(532), 1998–2010. doi:[10.1080/01621459.2019.1677240](https://doi.org/10.1080/01621459.2019.1677240).
- Parzer R, Filzmoser P, Vana-Gür L (2024a). “Data-Driven Random Projection and Screening for High-Dimensional Generalized Linear Models.” *Technical Report 2410.00971*, arXiv.org E-Print Archive. doi:[10.48550/arXiv.2410.00971](https://doi.org/10.48550/arXiv.2410.00971).
- Parzer R, Filzmoser P, Vana-Gür L (2024b). “Sparse Data-Driven Random Projection in Regression for High-Dimensional Data.” *Technical Report 2312.00130*, arXiv.org E-Print Archive. doi:[10.48550/arXiv.2312.00130](https://doi.org/10.48550/arXiv.2312.00130).
- Pedregosa F, Varoquaux G, Gramfort A, Michel V, Thirion B, Grisel O, Blondel M, Prettenhofer P, Weiss R, Dubourg V, *et al.* (2011). “**scikit-learn**: Machine Learning in Python.” *Journal of Machine Learning Research*, **12**(Oct), 2825–2830. URL <https://www.jmlr.org/papers/v12/pedregosa11a.html>.
- Raymaekers J, Rousseeuw P (2023). **cellWise**: *Analyzing Data with Cellwise Outliers*. doi:[10.32614/CRAN.package.cellWise](https://doi.org/10.32614/CRAN.package.cellWise). R package version 2.5.3.
- R Core Team (2024). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.

- Ryder N, Karnin Z, Liberty E (2019). “Asymmetric Random Projections.” *Technical Report 1906.09489*, arXiv.org E-Print Archive. doi:[10.48550/arXiv.1906.09489](https://doi.org/10.48550/arXiv.1906.09489).
- Saldana DF, Feng Y (2018). “SIS: An R Package for Sure Independence Screening in Ultrahigh-Dimensional Statistical Models.” *Journal of Statistical Software*, **83**(2), 1–25. doi:[10.18637/jss.v083.i02](https://doi.org/10.18637/jss.v083.i02).
- Siddharth R, Aghila G (2020). “RandPro – A Practical Implementation of Random Projection-Based Feature Extraction for High Dimensional Multivariate Data Analysis in R.” *SoftwareX*, **12**, 100629. ISSN 2352-7110. doi:[10.1016/j.softx.2020.100629](https://doi.org/10.1016/j.softx.2020.100629).
- Thanei GA, Heinze C, Meinshausen N (2017). *Random Projections for Large-Scale Regression*, pp. 51–68. Springer International Publishing, Cham. doi:[10.1007/978-3-319-41573-4_3](https://doi.org/10.1007/978-3-319-41573-4_3).
- Tian Y, Feng Y (2021). **RaSEn**: *Random Subspace Ensemble Classification and Variable Screening*. doi:[10.32614/CRAN.package.RaSEn](https://doi.org/10.32614/CRAN.package.RaSEn). R package version 3.0.0.
- van Rossum G, et al. (2011). *Python Programming Language*. URL <http://www.python.org>.
- Wang X, Leng C (2016). “High-dimensional Ordinary Least-squares Projection for Screening Variables.” *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, **78**, 589–611. doi:[10.1111/rssb.12127](https://doi.org/10.1111/rssb.12127).
- Wang X, Pan W, Hu W, Tian Y, Zhang H (2015). “Conditional Distance Correlation.” *Journal of the American Statistical Association*, **110**(512), 1726–1734. doi:[10.1080/01621459.2014.993081](https://doi.org/10.1080/01621459.2014.993081).
- Wanjun Liu Yuan Ke JL, Li R (2022). “Model-Free Feature Screening and FDR Control With Knockoff Features.” *Journal of the American Statistical Association*, **117**(537), 428–443. doi:[10.1080/01621459.2020.1783274](https://doi.org/10.1080/01621459.2020.1783274).
- Wickham H (2016). **ggplot2**: *Elegant Graphics for Data Analysis*. Springer-Verlag New York. ISBN 978-3-319-24277-4. URL <https://ggplot2.tidyverse.org>.
- Xu C, Chen J (2014). “The Sparse MLE for Ultrahigh-Dimensional Feature Screening.” *Journal of the American Statistical Association*, **109**(507), 1257–1269. doi:[10.1080/01621459.2013.879531](https://doi.org/10.1080/01621459.2013.879531).
- Zhu LP, Li L, Li R, Zhu LX (2011). “Model-Free Feature Screening for Ultrahigh-Dimensional Data.” *Journal of the American Statistical Association*, **106**(496), 1464–1475. doi:[10.1198/jasa.2011.tm10563](https://doi.org/10.1198/jasa.2011.tm10563).

Affiliation:

Roman Parzer

Computational Statistics (CSTAT), Institute of Statistics and Mathematical Methods in Economics

Wiedner Hauptstraße 8-10

Vienna Austria

E-mail: romanparzer1@gmail.com

Laura Vana Gür

Computational Statistics (CSTAT), Institute of Statistics and Mathematical Methods in Economics

Wiedner Hauptstraße 8-10

Vienna Austria

Peter Filzmoser

Computational Statistics (CSTAT), Institute of Statistics and Mathematical Methods in Economics

Wiedner Hauptstraße 8-10

Vienna Austria